



# Containers and algorithms — session 5

Tristan Brindle

# Feedback



- We'd love to hear from you!
- The easiest way is via the *cpplang* channel on Slack — we have our own chatroom, *#cpplondonuni*
- Go to <https://cpplang.now.sh/> for an “invitation”

# Last week



- Iterator categories
- An overview of standard containers

# This week

- Finishing the container overview
- Algorithmic complexity
- Intro to STL algorithms

# Standard containers overview



- The C++ standard library provides several container classes which can (and should!) be used when writing our own programs
- These can broadly be divided into four categories:
  - Sequence containers: elements are stored in the order they are added
  - Associative containers: keys are *sorted* for fast lookup
  - Unordered associative containers: keys are *hashed* for fast lookup
  - Container adaptors: provide a modified interface for specific tasks

# Standard containers overview



- All standard containers are *class templates*
- This means that when you create an *instance* of a container to hold a particular type, it is specialised *just for that type*
- This means that a `std::vector<int>` is not the same type as `std::vector<double>`, and `std::list<std::string>` is not the same type as `std::list<bool>`
- Some containers have parameters which can be used to specialise the behaviour — for example the unordered containers can use a customised hash function
- All the standard containers have an `Allocator` template parameter which can be used to optimise memory allocations in advanced use cases

# Standard containers

## overview



- The standard containers are designed to have a consistent programmer interface
- For example, `std::vector`, `std::list` and `std::deque` all have a `push_back()` member function which appends an element to the end of the sequence
- In particular, all the containers have `begin()` and `end()` member functions which return *iterators*
- They also have `cbegin()` and `cend()` functions which return *const iterators* — these provide read-only access to the container
- See <https://en.cppreference.com/w/cpp/container> for complete details

# Sequence containers



- `std::array`: Fixed-size random-access array
- `std::vector`: Dynamically-sized random-access array
- `std::list`: Doubly-linked list
- `std::forward_list`: Singly-linked list
- `std::deque`: Double-ended queue
- `std::string`: Dynamically-sized random-access array of characters



# std::vector



- std::vector is the standard's version of a dynamically-sized, random-access array
- Properties:
  - Element access is constant time
  - Adding elements to the end of a vector is *amortised constant time*
  - Adding elements anywhere else is *linear* in the number of elements in the vector
  - Iterators are random-access (guaranteed contiguous in C++17)
- std::vector should be your default, go-to container for most uses
- **When in doubt, use vector!**

# std::list



- `std::list` is a doubly-linked list
- Properties:
  - Adding/removing an element anywhere in the list is constant time
  - Accessing a particular element is linear in the number of elements in the list
  - Accessing the `size()` of the list is constant time in C++11
  - Iterators are *bidirectional*
- Note that `std::list` is *node-based*, meaning adding an element requires a dynamic allocation
- Guideline: prefer `std::list` to `std::vector` only when frequently inserting and removing elements from the middle of a *large* sequence

# std::deque

- `std::deque` is a *double-ended queue*
- Properties:
  - Element access is constant time
  - Insertion or removal at the end **or the beginning** of a deque is constant time
  - Insertion or removal elsewhere is linear in the size of the deque
  - Iterators are random access
- `std::deque` offers the same complexity guarantees as `std::vector`, plus an extra guarantee regarding insertion at the start
- But there's no such thing as a free lunch: deque's operations have a larger constant factor, and there is likely to be extra memory overhead

# Associative containers

- `std::map`: A collection of unique key-value pairs, sorted by keys
- `std::set`: A sorted collection of unique keys
- `std::multimap`: A collection of (possibly non-unique) key-value pairs, sorted by keys
- `std::multiset`: A sorted collection of (possibly non-unique) keys

# std::map

- `std::map` is a sorted container of unique key-value pairs
- Properties:
  - Element access and insertion is  $O(\log n)$
  - Iterators are *bidirectional*
  - Iteration order is well-defined
- Typically implemented as a *red-black tree*
- By default, keys are sorted using operator<. This can be customised using map's Compare template parameter
- `std::map` should be your default, go-to associative array type for most purposes

# `std::set`



- A `std::set` is a sorted container of unique keys
- Think of it as a map without the values!
- `std::set` can be useful for ensuring that you have a collection of unique elements
- However, you can often achieve better performance by using a `std::vector` and keeping it sorted yourself

# Unordered associative containers

- `std::unordered_map`
- `std::unordered_set`
- `std::unordered_multimap`
- `std::unordered_multiset`

# std::unordered\_map



- `std::unordered_map` is the standard library's version of a hash table
- Properties:
  - Element access and insertion:  $O(1)$  average,  $O(n)$  worst-case
  - Iterators are forward only
  - Iteration order is unspecified
- Compared with `std::map`, `unordered_map` offers constant-time lookup on average
- Keys are compared using `std::hash` by default. This is defined for all built-in types, but you need to specialise it for your own types
- As with all hash tables, performance is highly dependent on the quality of the hash function



# Container adaptors



- The container adaptors in the STL take an underlying sequence container and wrap it in a new, more specialised, more restrictive API
- The adapted classes are not containers themselves (they have no `begin()` or `end()`), so in practise are rarely used
- The container adaptors are:
  - `std::queue`: Adapts a container (such as vector) into a FIFO queue
  - `std::stack`: Adapts a container into a LIFO stack
  - `std::priority_queue`: Adapts a container so as to provide constant-time access to the largest element

# Exercise



- Clone the repository at

[https://github.com/CPPLondonUni/  
stl\\_week4\\_class\\_exercise](https://github.com/CPPLondonUni/stl_week4_class_exercise)

- Please complete exercise 1 in the README file

# Solution



- [https://github.com/CPPLondonUni/stl\\_week4\\_class\\_exercise/tree/ex1\\_solution](https://github.com/CPPLondonUni/stl_week4_class_exercise/tree/ex1_solution)

# Online resources



- <https://isocpp.org/get-started>
- [cppreference.com](http://cppreference.com) — The bible, but aimed at experts
- [cplusplus.com](http://cplusplus.com) — Another reference site, also has a tutorial section
- [learncpp.com](http://learncpp.com) — Free online tutorial, very up-to-date
- <https://www.pluralsight.com/authors/kate-gregory> - Comprehensive set of courses from an experienced C++ trainer (free trial)
- [reddit.com/r/cpp\\_questions](https://reddit.com/r/cpp_questions)
- Cpplang Slack channel — <https://cpplang.now.sh/> for an “invite”
- StackOverflow (but...)