



Containers and algorithms — session 6

Tristan Brindle

Feedback



- We'd love to hear from you!
- The easiest way is via the *cpplang* channel on Slack — we have our own chatroom, *#cpplondonuni*
- Go to <https://cpplang.now.sh/> for an “invitation”

Last week



- Further STL container overview

This week

- Algorithmic complexity
- Intro to standard algorithms

Last week's homework



Last week's homework



- Clone the repository at

[https://github.com/CPPLondonUni/
stl_week4_class_exercise](https://github.com/CPPLondonUni/stl_week4_class_exercise)

Solution



- https://github.com/CPPLondonUni/stl_week4_class_exercise/tree/ex1_solution

**Any questions before
we move on?**

Algorithmic complexity

- In computer science, the *complexity* of an algorithm refers to the amount of resources required to execute it
- We are mostly interested in *time complexity*, though *space complexity* is also important for some problems
- Algorithmic analysis is a complicated and deeply theoretical subject. We will only touch on the very very basics today
- For a more rigorous mathematical treatment, consult a computer science textbook 😊

Big-O notation

- Most of the time, we are interested in how an algorithm's performance *scales* as we increase the problem size
- We can express this using “big-O” notation
- Roughly, given a (mathematical) function $f(n)$, its behaviour as n increases is dominated by the *fastest-growing term*
- We call this the *order* of the function, denoted $O(x)$

Big-O notation

- For example, if a given algorithm on n elements takes


$$f(n) = 4n^2 + 3n + 8$$

operations, then as n grows very large, $f(n)$ will be dominated by the n^2 term

- We say that this algorithm has order n^2 , or $O(n^2)$

Big-O notation

- If an algorithm does not depend on the number of elements, we say it operates in *constant time*, written $O(1)$
- An $O(n)$ algorithm is said to operate in linear time
- Other common complexity classes:
 - $O(\log n)$: logarithmic time
 - $O(n^2)$: quadratic time
 - $O(n^k)$: polynomial time
 - $O(2^n)$: exponential time



By Cmglee - Own work, CC BY-SA 4.0, <https://commons.wikimedia.org/w/index.php?curid=50321072>

Big-O is not everything!

- The theoretical complexity is important in choosing an appropriate algorithm
- However, “big-O” masks constant factors and lower-order terms, which can be more important for real-world problem sizes
- In particular, caching and modern CPU optimisations can have surprising effects
- When in doubt — **measure!**

Complexity and the STL



- For containers and algorithms, the C++ standard generally does not specify exactly the implementation that should be used
- Rather, it specifies the **algorithmic complexity** that operations must have
- For example, `std::vector`'s operator[] must operate in constant time. `std::sort()` must have $O(n \cdot \log n)$ complexity.

**Any questions before
we move on?**

Standard algorithms



- The standard library contains over 90 different algorithms!
- These are mostly defined in header `<algorithm>`, with some extras like `std::accumulate()` defined in header `<numeric>`
- These range from the very basic (e.g. `std::count()`) to the very specialised (e.g. `std::sort()`)
- General advice: use the standard algorithms whenever you can!
- Further: if you find yourself writing a non-trivial for loop, see if you can abstract out the operation into a self-contained algorithm
- See Sean Parent's fantastic C++ Seasoning talk for inspiration

Standard algorithms



- The standard library algorithms operate on pairs of iterators, called a *range*
- The first iterator of the pair denotes the *starting element*
- The last iterator of the pair denotes *one past the last element*
- Passing `container.begin()` and `container.end()` will operate on the whole container, but it is also possible to operate on a subset.

Example

// Imaginary function returning a vector of integers

```
auto vec = get_some_ints();
```

```
auto num_zeros = std::count(vec.begin(), vec.end(), 0);
```

```
std::cout << "There are " << num_zeros << " zeroes in the whole vector\n";
```

```
std::fill(vec.begin(), vec.begin() + 10, 42);
```

// Fills the first 10 elements of the vector with the value 42

Anatomy of a standard algorithm

- The standard algorithms are implemented as *function templates*
- A template is like a *blueprint*, telling the compiler *how* to create the function
- When we call a function template, the compiler creates a *specialised* version for the argument types we supply
- This allows for unbeatable efficiency!

Anatomy of a standard algorithm



// count() implementation

```
template <typename Iter, typename T>
std::size_t count(Iter first, Iter last, const T& value)
{
    std::size_t counter = 0;
    for ( ; first != last; ++first) {
        if (*first == value) {
            ++counter;
        }
    }
    return counter;
}
```

// Calling count()

```
const std::vector<int> ints{1, 2, 3, 4, 1};
```

```
const auto num_ones = count(ints.begin(), ints.end(), 1);
```

// Compiler creates a version of count() with T = int, Iter = std::vector<int>::iterator

```
const std::list<float> floats{1.0f, 2.0f, 3.0f};
```

```
const auto num_twos = count(floats.begin(), floats.end(), 2.0);
```

// Compiler creates a version of count() with T = double, Iter = std::list<float>

Predicates

- Many standard algorithms allow you to pass *predicates* which dictate their behaviour
- A predicate is just a *callable* (a function, function object or lambda) which returns a `bool`
- For now we'll just be using regular functions
- Some algorithms require *unary predicates* taking one argument (usually of the iterator's `value_type`)
- Other algorithms require a *binary predicate*, comparing two elements of the range

Using predicates



```
struct Person {
    std::string first_name;
    std::string last_name;
};

bool compare_person(const Person& a, const Person& b)
{
    if (a.last_name == b.last_name) {
        return a.first_name < b.first_name;
    }
    return a.last_name < b.last_name;
}

std::vector<Person> people;
/* ...fill people vector... */

// Get iterator to the first person, sorted alphabetically
auto iter = std::min_element(people.begin(), people.end(), compare_person);

// Dereference the iterator to get a reference to the Person instance
Person& first_person = *iter;
```

count_if()



```
// count_if() implementation
template <typename Iter, typename Pred>
std::size_t count_if(Iter first, Iter last, Pred pred)
{
    std::size_t counter = 0;
    for (; first != last; ++first) {
        if (pred(*first)) {
            ++counter;
        }
    }
    return counter;
}

// Calling count_if
bool starts_with_T(const std::string& str) { return !str.empty() && str[0] == 'T'; }

const std::vector<std::string> people{"Tom", "Oli", "Tristan"};
const auto num_Ts = count_if(people.begin(), people.end(), starts_with_T);

// could also use a lambda

const auto num_short_names = count_if(people.begin(), people.end(),
                                       [](const auto& str) {
                                           return str.size() <= 3;
                                       });
```


Exercise



- Clone the repository at
- https://github.com/CPPLondonUni/algorithms_exercise
- Your task is to implement some standard library algorithms given the specification
- Tests are included

Online resources



- <https://isocpp.org/get-started>
- cppreference.com — The bible, but aimed at experts
- cplusplus.com — Another reference site, also has a tutorial section
- learncpp.com — Free online tutorial, very up-to-date
- <https://www.pluralsight.com/authors/kate-gregory> - Comprehensive set of courses from an experienced C++ trainer (free trial)
- reddit.com/r/cpp_questions
- Cpplang Slack channel — <https://cpplang.now.sh/> for an “invite”
- StackOverflow (but...)